

# notebook

May 21, 2020

## IB031 Project

- [Source](#)
- [Testing csv](#)
- [Training csv](#)

In this project we will implement several models to predict credibility of a loan applicant. We will use the Loan Prediction Problem Dataset from Kaggle. The structure of the project is as follows:

1. Section ??
2. Section ??
3. Section ??
4. Section ??
5. Section ??
6. Section ??
7. Section ??
8. Section ??
9. Section ??

### 0.1 1. Exploratory Analysis

The dataset consists of basic information about applicants. There are 13 columns with about 1000 rows. First of all, we will deal with the data types.

```
[2]: import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

sns.set()
```

Since pandas loads categorical data types as ‘*object*’ dtype by default, we want to convert them back to category after loading the csv. We are also dropping the “Loan\_Status” column, which will later provide labels for classification. Unfortunately, this particular dataset does not contain a set of testing labels, therefore we were forced to split an already sparse training set into 2 parts.

```
[3]: from sklearn.model_selection import train_test_split

#X_test = pd.read_csv("./test_Y3wMUE5_7gLdaTN.csv", index_col="Loan_ID")
```

```
dataset = pd.read_csv("./train_u6lujuX_CVtuZ9i.csv", index_col="Loan_ID")

for col in dataset.columns:
    if dataset[col].dtype == 'object':
        dataset[col] = dataset[col].astype('category')

X, y = dataset.drop(["Loan_Status"], axis=1), dataset["Loan_Status"][:]
y = y.map({"Y": True, "N": False})

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2)
```

[4]: X.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 614 entries, LP001002 to LP002990
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                 601 non-null   category
1   Married                611 non-null   category
2   Dependents             599 non-null   category
3   Education              614 non-null   category
4   Self_Employed          582 non-null   category
5   ApplicantIncome        614 non-null   int64
6   CoapplicantIncome      614 non-null   float64
7   LoanAmount             592 non-null   float64
8   Loan_Amount_Term       600 non-null   float64
9   Credit_History         564 non-null   float64
10  Property_Area          614 non-null   category
dtypes: category(6), float64(4), int64(1)
memory usage: 33.0+ KB
```

We have dropped the 'Loan\_Status' column, which will provide labels during classification

[5]: X.shape

[5]: (614, 11)

[6]: y.shape

[6]: (614,)

The dataset had been split into training and testing subsets in around 1.67:1 ratio.

[7]: X.head()

```
[7]:      Gender Married Dependents      Education Self_Employed \
Loan_ID
```

LP001002	Male	No	0	Graduate	No
LP001003	Male	Yes	1	Graduate	No
LP001005	Male	Yes	0	Graduate	Yes
LP001006	Male	Yes	0	Not Graduate	No
LP001008	Male	No	0	Graduate	No

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
Loan_ID				
LP001002	5849	0.0	NaN	360.0
LP001003	4583	1508.0	128.0	360.0
LP001005	3000	0.0	66.0	360.0
LP001006	2583	2358.0	120.0	360.0
LP001008	6000	0.0	141.0	360.0

	Credit_History	Property_Area
Loan_ID		
LP001002	1.0	Urban
LP001003	1.0	Rural
LP001005	1.0	Urban
LP001006	1.0	Urban
LP001008	1.0	Urban

```
[8]: X.isnull().sum()
```

```
[8]: Gender          13
      Married         3
      Dependents     15
      Education       0
      Self_Employed  32
      ApplicantIncome  0
      CoapplicantIncome 0
      LoanAmount      22
      Loan_Amount_Term 14
      Credit_History   50
      Property_Area    0
      dtype: int64
```

As we can see, there are a lot of null values which will need to be imputed. Since there is no column with more missing data than present, we do not need to drop any.

```
[9]: X.describe()
```

```
[9]:      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
count      614.000000      614.000000    592.000000      600.00000
mean      5403.459283      1621.245798    146.412162      342.00000
std       6109.041673      2926.248369     85.587325       65.12041
min       150.000000         0.000000     9.000000      12.00000
```

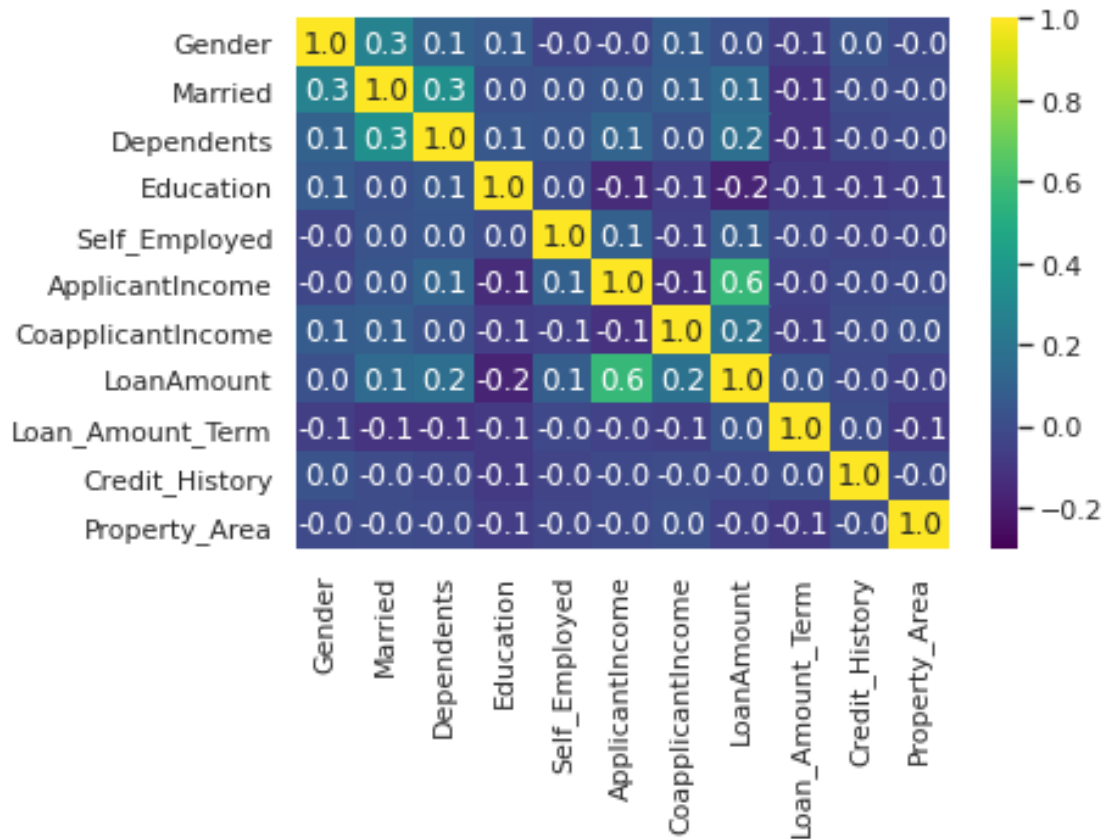
25%	2877.500000	0.000000	100.000000	360.000000
50%	3812.500000	1188.500000	128.000000	360.000000
75%	5795.000000	2297.250000	168.000000	360.000000
max	81000.000000	41667.000000	700.000000	480.000000

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

Here we have some essential statistics about our dataset. For instance, we can see, that the most frequent length of loan term is about 1 year, with maximum being 1.5 year and minimum just 12 days. Furthermore, coapplicants have much lower income than applicants.

```
[10]: corr_X = X[:]
for col in {"Property_Area", "Dependents", "Gender", "Married", "Education",
           ↪ "Self_Employed"}:
    corr_X[col] = corr_X[col].cat.codes

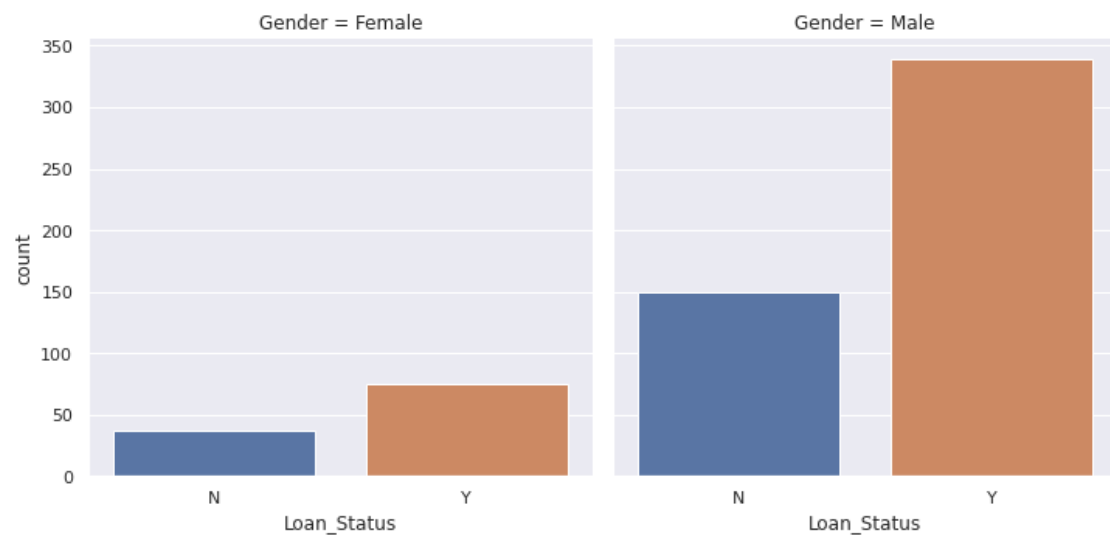
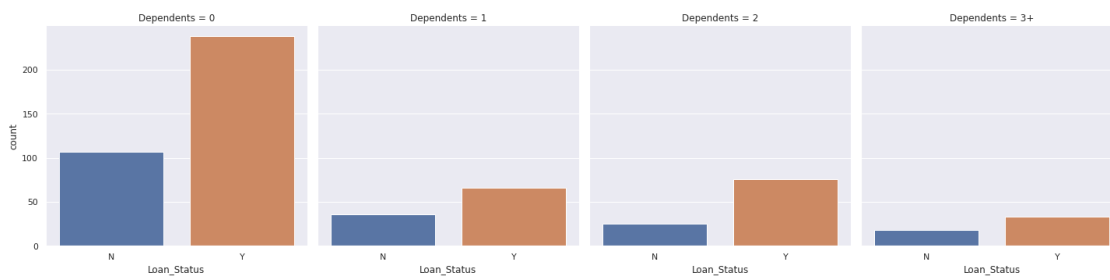
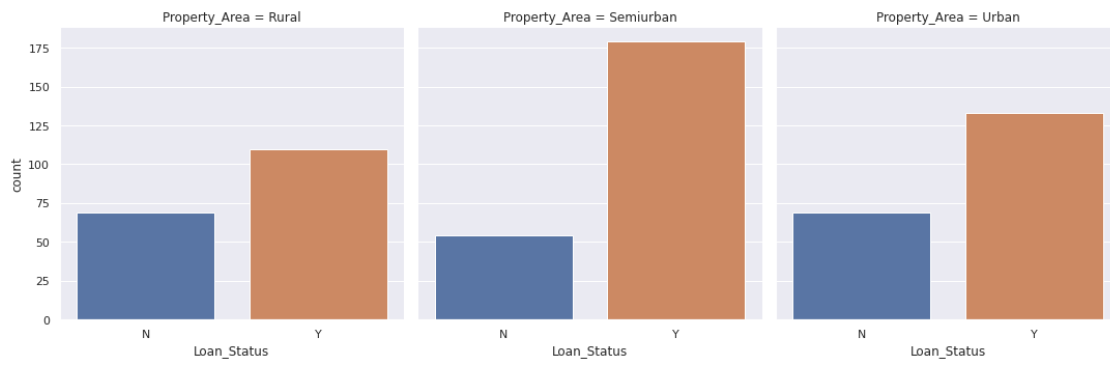
sns.heatmap(corr_X.corr(), annot=True, vmin=-0.3, cmap="viridis", fmt=".1f")
plt.show()
```




We also computed and plotted correlation matrix of features in the dataset. There is only one pair of features, loan amount and applicant income, which can be considered to be mildly positively correlated. Other features, have nearly no correlation among them. This results is quite suprising, as one would expect higher correlation of features.

```
[11]: sns.catplot("Loan_Status", col="Property_Area", data=dataset, kind="count")
sns.catplot("Loan_Status", col="Dependents", data=dataset, kind="count")
sns.catplot("Loan_Status", col="Gender", data=dataset, kind="count")
sns.catplot("Loan_Status", col="Married", data=dataset, kind="count")
sns.catplot("Loan_Status", col="Education", data=dataset, kind="count")
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x7f6595abf9b0>
```







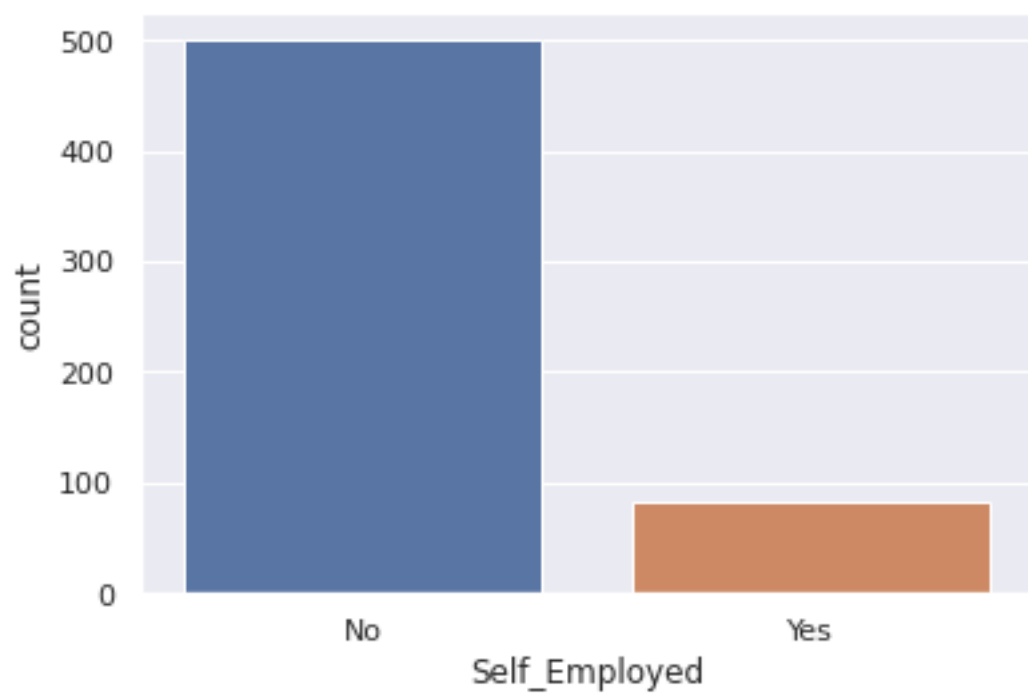
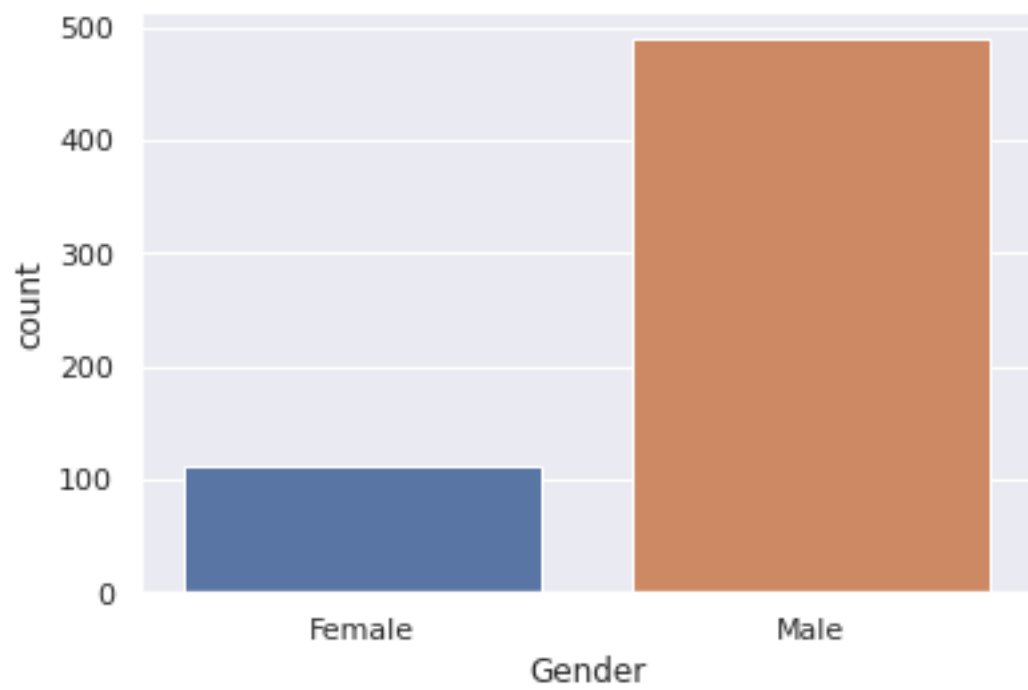
We can look at the amount of distinct values of selected features.

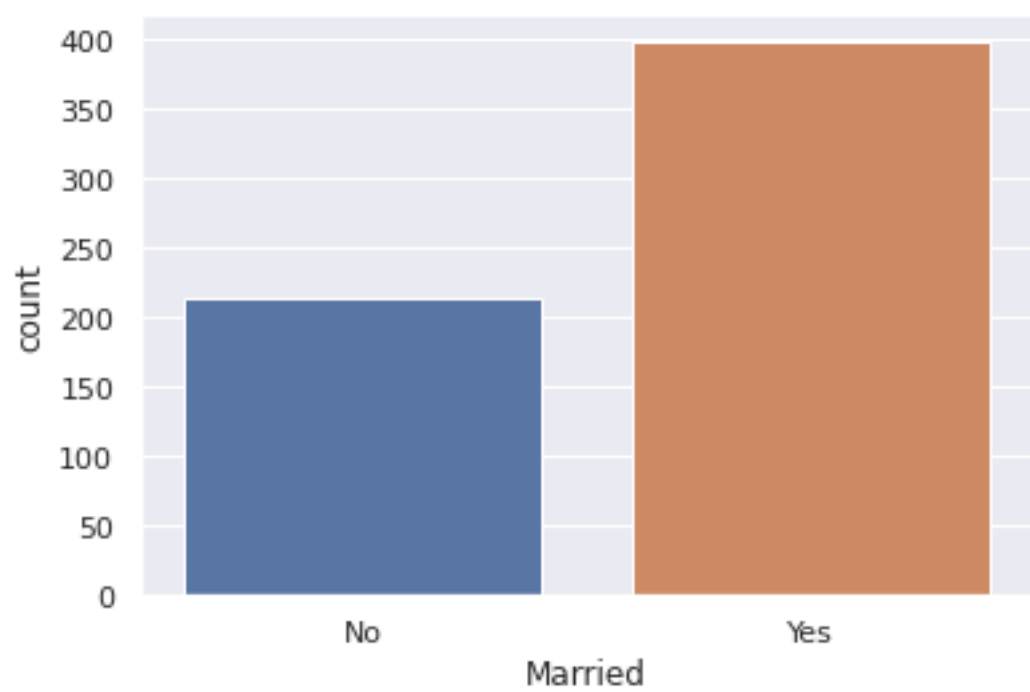
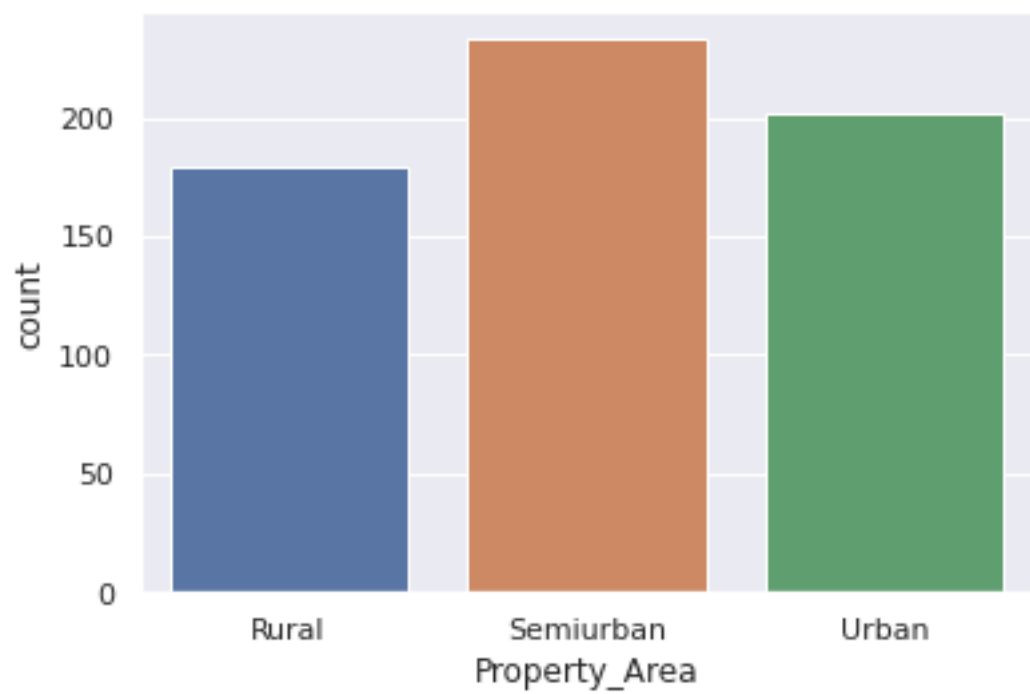
```
[36]: for col in {"Property_Area", "Dependents", "Gender", "Married", "Education",
    ↪ "Self_Employed"}:
    sns.countplot(X[col])
    plt.show()


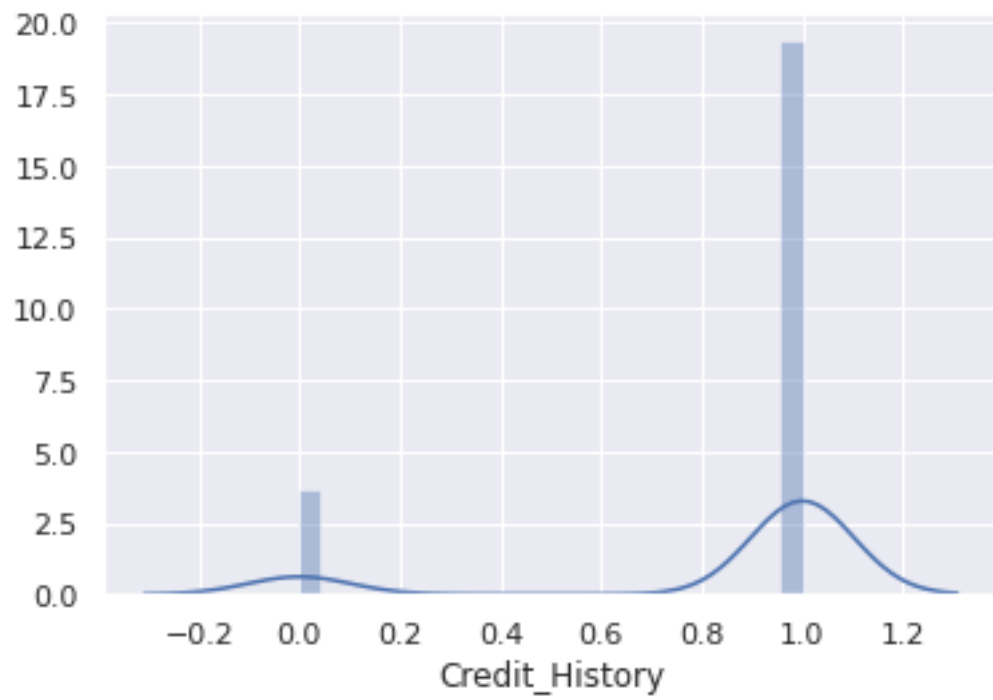
for col in {"ApplicantIncome", "CoapplicantIncome", "Credit_History",
    ↪ "LoanAmount", "Loan_Amount_Term"}:
    sns.distplot(X[col])
    plt.show()
```

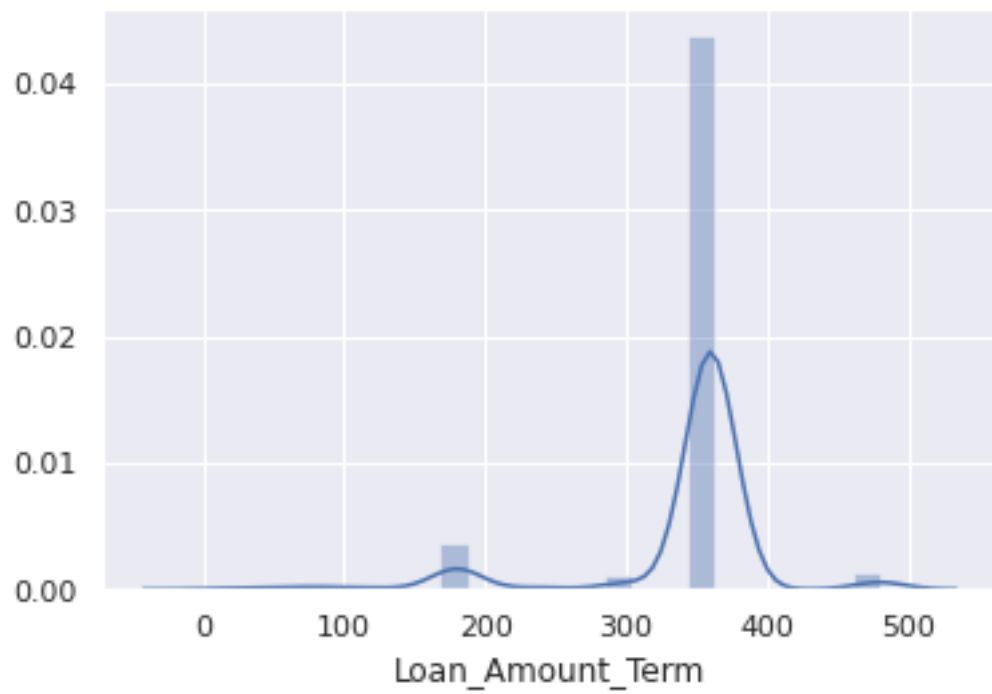



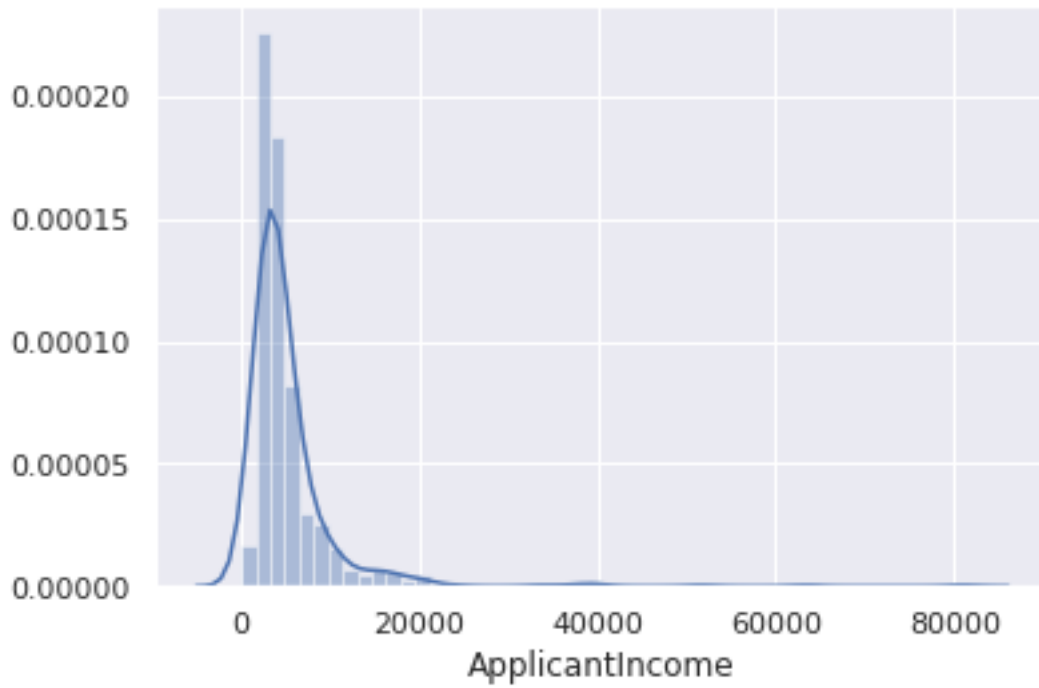













Most of the features are scattered along the x-axis. The only exception is Loan Amount, which roughly follows the Chi-squared distribution.

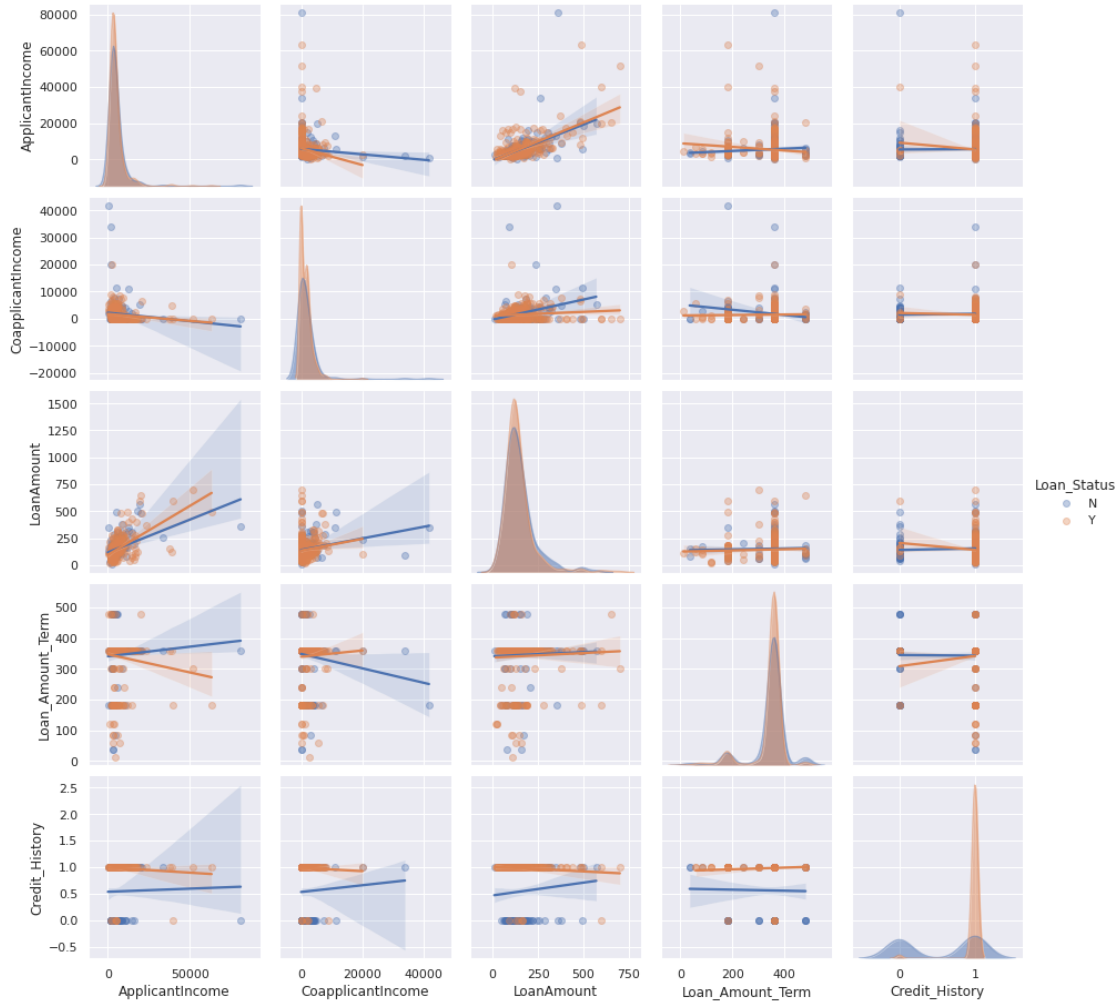
```
[13]: sns.countplot(y)
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f6594e9aa58>
```



```
[14]: sns.pairplot(dataset, hue="Loan_Status", kind="reg", diag_kws={"alpha": 0.5},  
      ↪ plot_kws={"scatter_kws": {"alpha": 0.35}})
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x7f6594e72ef0>
```



Credit history seems to have the highest impact on deciding whether the loan should be approved or not. Clients with low credit score seem more likely to be denied.

There are some cases where ApplicantIncome and CoapplicantIncome are very low, their loan request gets approved, whereas in a few cases the CoapplicantIncome is very high and ApplicantIncome is very low, the request gets denied.

## 0.2 2. Data preprocessing

```
[15]: y.isna().sum()
```

```
[15]: 0
```

Good, there are no N/A labels. We do not need to drop any rows. Besides that, dropping a significant part of the dataset could have misrepresenting effects.

Firstly, we will impute the missing values with `SimpleImputer` and strategy `most_frequent`. As `SimpleImputer` returns an array, we will transform it back to a pandas dataframe. Another step

is encoding the categorical features using `OneHotEncoder`. Lastly, we will scale the features with `StandardScaler`. Their mean will therefore be 0 and variance 1.

The labels need significantly less preprocessing. Encoding their boolean values is sufficient.

```
[16]: from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder, StandardScaler
      from sklearn.pipeline import make_pipeline
      from sklearn.compose import make_column_transformer
      from sklearn.impute import SimpleImputer

      class NpToDf:
          columns = []

          def __init__(self, columns=None):
              self.columns = columns

          def fit(self, X, *args, **kwargs):
              return X

          def fit_transform(self, X, *args, **kwargs):
              return self.transform(X)

          def transform(self, X, *args, **kwargs):
              return pd.DataFrame(data=X, columns=self.columns)

      pipe_X = make_pipeline(
          SimpleImputer(missing_values=np.nan, strategy="most_frequent"),
          NpToDf(X_train.columns),
          make_column_transformer(
              (OneHotEncoder(), ['Gender', 'Married', 'Education', 'Self_Employed',
                               ↪ 'Property_Area', 'Dependents']),
              remainder=StandardScaler()),
          NpToDf(),
      )

      pipe_y = make_pipeline(
          OrdinalEncoder()
      )
```

```
[17]: pipe_X.fit_transform(X_train)
      train_X = pipe_X.transform(X_train)
      test_X = pipe_X.transform(X_test)

      train_y = pipe_y.fit_transform(y_train.to_numpy().reshape(-1, 1)).reshape(-1)
      test_y = pipe_y.fit_transform(y_test.to_numpy().reshape(-1, 1)).reshape(-1)
```

Some helper functions for training and evaluation of our models.



```
[18]: from sklearn.metrics import mean_squared_error, f1_score
from sklearn.model_selection import cross_val_score
import matplotlib.pyplot as plt

def evaluate(clf, X_test, y_test):
    y_pred = clf.predict(X_test)
    scores = cross_val_score(clf, X_test, y_test, cv=10)

    print(f"RMSE: {mean_squared_error(y_test, y_pred, squared = False):.4f}")
    print(f"Accuracy: {scores.mean():.3f} ± {scores.std() * 2:.3f}")
    print(f"F1 Score: %.2f" % f1_score(y_test, y_pred, average='weighted'))

[19]: from sklearn.metrics import plot_roc_curve

def roc(clf, test_X, test_y):
    plot_roc_curve(clf, test_X, test_y)
    plt.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance',
    ↪alpha=.8)
    plt.legend()

[20]: from sklearn.metrics import plot_precision_recall_curve as prc
from sklearn.metrics import precision_recall_curve

[21]: from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV

def confusion(clf, X_test, y_test):
    plot_confusion_matrix(clf, X_test, y_test, cmap=plt.cm.Blues)

[22]: def get_gscv(clf, param_grid, verbose=1, **kwargs):
    gs = GridSearchCV(clf, param_grid, verbose=verbose, cv=kwargs.get("cv", 3),
    ↪n_jobs=kwargs.get("workers", -2))
    gs.fit(train_X, train_y, **kwargs)
    score = gs.score(test_X, test_y)
    print(f"Best parameters: {gs.best_params_}, with F1 score of {score:.2f}")
    return gs.best_estimator_
```

In the following 4 sections, we will always start by running a grid search tuning hyperparameters of our models. There is a wrapper for keras sequential models, which we will use for grid searching. All models will be trained on the same training set. We chose these evaluation metrics - **RMSE** - measures error of the predictions compared to actual values, the lower the better. - **Accuracy** computed with cross validation score - is the proportion of correct predictions (both true positives and true negatives) among the total number of cases examined, the higher the better - **F1 score** - is weighted average of precision and recall, the best value is 1, the worst is 0 - **Receiver Operating Characteristic curve** - measures the ability of a model to distinguish between classes - **Precision**

**Recall curve** - shows the tradeoff between precision and recall for different threshold - **Confusion matrix** - shows the number of True Positive (TP), False Negative (FN), True Negative (TN), False Positive (FP) classifications.

### 0.3 3. Naive baseline model

This is a simple classifier, which chooses the class based on training set class distribution.

It is very basic and is affected by the chosen train/test split a lot.

```
[23]: from sklearn.dummy import DummyClassifier

dummy_clf = DummyClassifier(strategy="stratified")

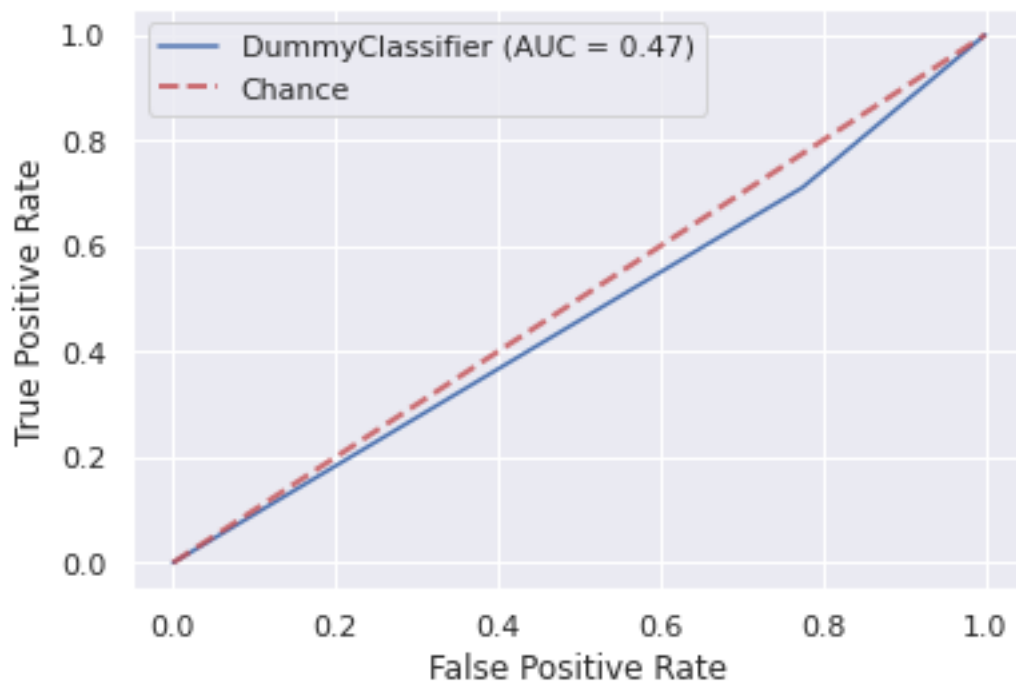
dummy_clf.fit(train_X, train_y)


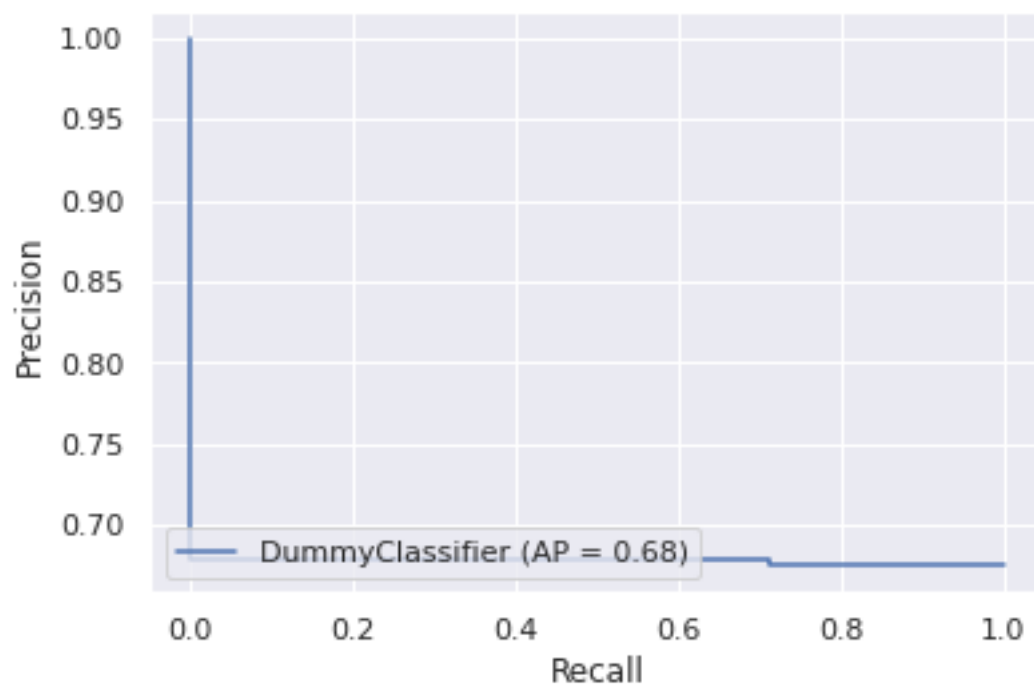
evaluate(dummy_clf, test_X, test_y)
roc(dummy_clf, test_X, test_y)
prc(dummy_clf, test_X, test_y)
confusion(dummy_clf, test_X, test_y)
plt.show()
```

RMSE: 0.6564

Accuracy: 0.610 ± 0.176

F1 Score: 0.57





## 0.4 4. Decision tree classifier

Decision trees are among the most used classification models. They iteratively split the dataset, until a tree conforming to given parameters has been constructed. In leaves they contain class labels. Internal nodes represent kind of a boolean test, usually a value of a sample's feature, according to which the algorithm chooses the respective edge on the way to leaves. The tests can also use entropy and information gain to choose the best edge. There are many to ways to construct a tree, therefore extensive hyperparameter tuning is suitable. Decision trees can also be pruned, either during construction or after it.

```
[38]: len(train_X.columns)
```

```
[38]: 20
```

```
[46]: from sklearn.tree import DecisionTreeClassifier

dtree_clf = DecisionTreeClassifier()
dtree_values = {"criterion": ["gini", "entropy"],
                "max_depth": [1, 2, 5, 10, 16, None],
                "max_leaf_nodes": [6, 8, 10, 12, 20, None],
                "max_features": ["auto", "sqrt", "log2", 2, 4, 8, 11, 15,
                                len(train_X.columns)]}

tree_clf = get_gscv(dtree_clf, dtree_values)
```

Fitting 3 folds for each of 648 candidates, totalling 1944 fits

[Parallel(n\_jobs=-2)]: Using backend LokyBackend with 3 concurrent workers.

[Parallel(n\_jobs=-2)]: Done 634 tasks | elapsed: 2.3s

Best parameters: {'criterion': 'entropy', 'max\_depth': 10, 'max\_features': 15, 'max\_leaf\_nodes': 8}, with F1 score of 0.71



[Parallel(n\_jobs=-2)]: Done 1944 out of 1944 | elapsed: 6.9s finished


```
[47]: evaluate(tree_clf, test_X, test_y)
roc(tree_clf, test_X, test_y)
prc(tree_clf, test_X, test_y)
confusion(tree_clf, test_X, test_y)
plt.show()
```

RMSE: 0.5410

Accuracy: 0.758 ± 0.231

F1 Score: 0.63





## 0.5 5. KNN classifier

Another popular classification algorithm, an example of instance-based learning or lazy learning. This time, all distances from a data point to other points are computed, and  $k$ -closest neighbours are chosen. Then, the class memberships of the  $k$ -closest members are considered, with the original data point taking a class label from the most occurring one among its  $k$ -closest neighbours. For the distance metrics, *Euclidean* or *Hamming* distances are usually used. There is a tradeoff in the number of  $k$ -closest neighbours. Smaller  $k$ , signifies the result of noise on classification, but makes the various classes more distinct and vice versa with higher  $k$ .

```
[26]: from sklearn.neighbors import KNeighborsClassifier

knn_clf = KNeighborsClassifier()
knn_values = {"n_neighbors": list(range(2, 16, 2)),
              "algorithm": ["auto", "ball_tree", "kd_tree", "brute"],
              "leaf_size": [10, 20, 30, 40, 50],
              "p": [1, 2],
              "weights": ["uniform", "distance"]}

knn_clf = get_gscv(knn_clf, knn_values)
```

Fitting 3 folds for each of 560 candidates, totalling 1680 fits

[Parallel(n\_jobs=-2)]: Using backend LokyBackend with 3 concurrent workers.

[Parallel(n\_jobs=-2)]: Done 250 tasks | elapsed: 2.6s

[Parallel(n\_jobs=-2)]: Done 1450 tasks | elapsed: 13.9s

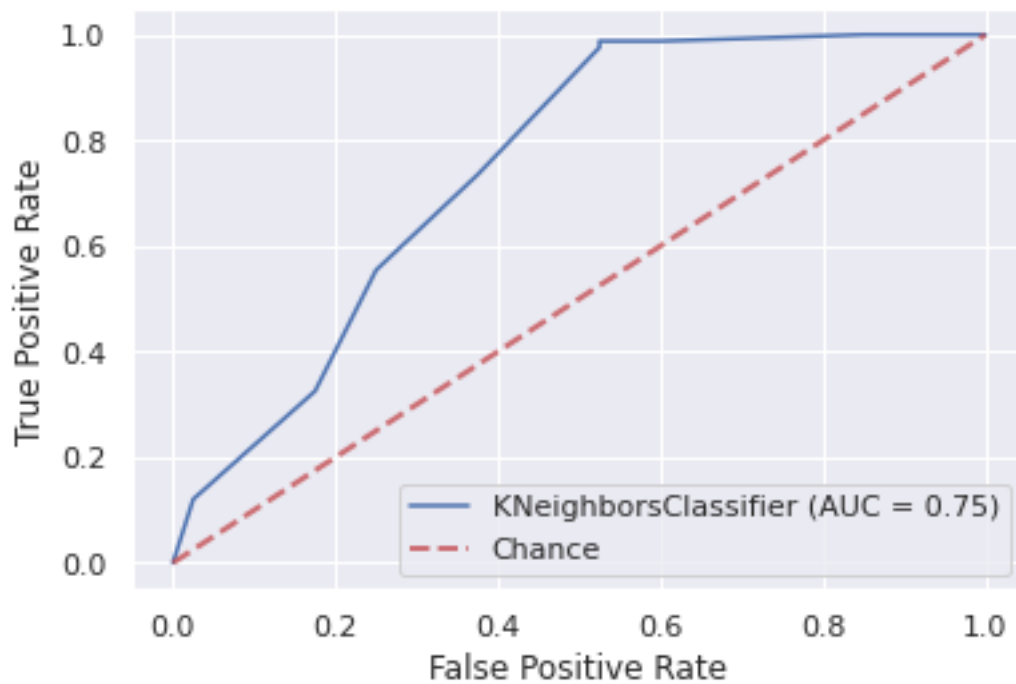
Best parameters: {'algorithm': 'auto', 'leaf\_size': 10, 'n\_neighbors': 14, 'p': 2, 'weights': 'uniform'}, with F1 score of 0.82  
[Parallel(n\_jobs=-2)]: Done 1680 out of 1680 | elapsed: 15.9s finished


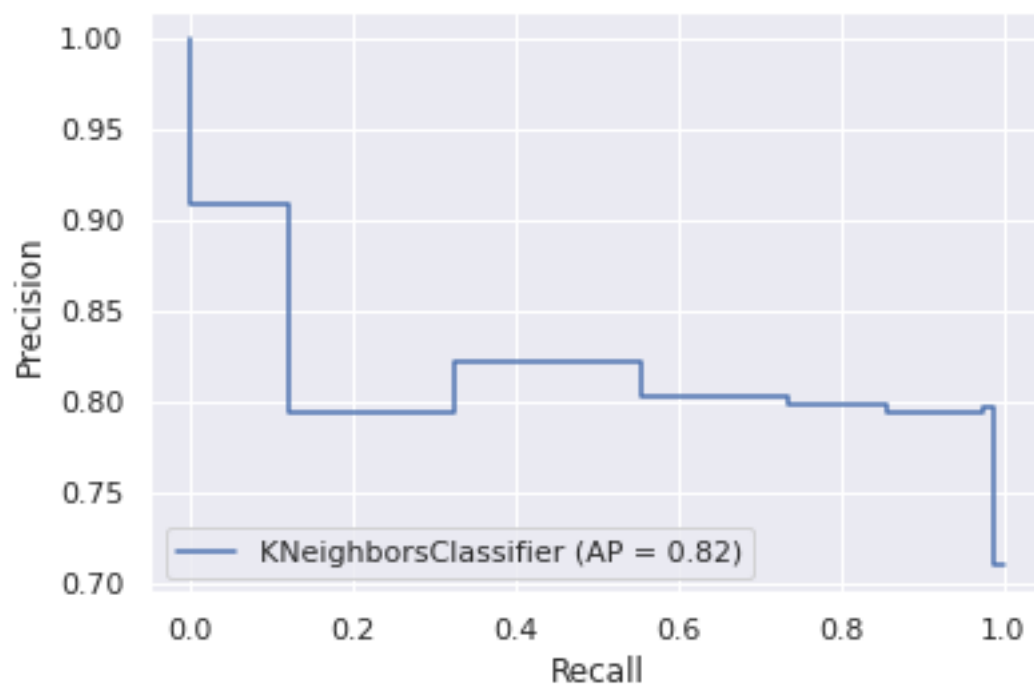
```
[27]: evaluate(knn_clf, test_X, test_y)
      roc(knn_clf, test_X, test_y)
      prc(knn_clf, test_X, test_y)
      confusion(knn_clf, test_X, test_y)
      plt.show()
```

RMSE: 0.4229

Accuracy: 0.804 ± 0.196

F1 Score: 0.80







## 0.6 6. Support Vector Machine

Support Vector Machines, abbr. SVM, is a supervised-learning algorithm used mainly for binary classification, although it is possible to use for multi-class classification by combining several SVMs. It creates hyperplanes in a multi-dimensional feature space, which are then used for generalization and classifications of data points. The best performing hyperplanes are those having the biggest maximum margin, i.e. the closest data points from both classes are as far as possible. In order to transform input data into a desired form, SVM uses so called kernel functions, which return the inner product between two points in a suitable feature space.

```
[28]: from sklearn.svm import SVC


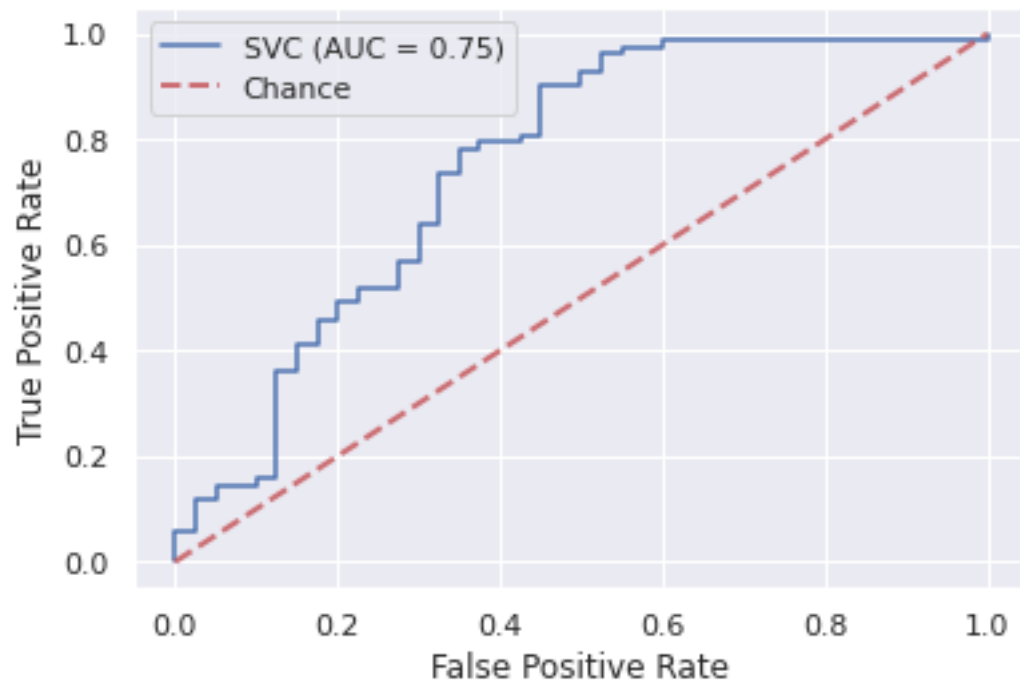
svc_clf = SVC()
svc_values = {"C": [.8, 1.0, 1.2, 1.5, 2.0],
              "kernel": ["linear", "poly", "rbf", "sigmoid"],
              "degree": [2, 3, 4, 5],
              "gamma": ["scale", "auto"],
              "coef0": [.0, .5, 1.0, 1.5]}

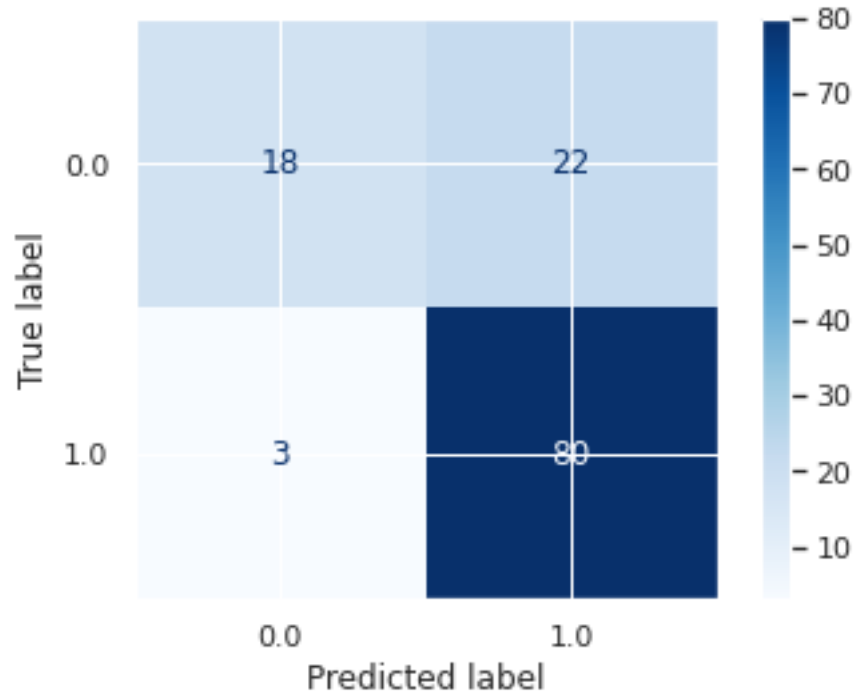
svm_clf = get_gscv(svc_clf, svc_values)
```

```
Fitting 3 folds for each of 640 candidates, totalling 1920 fits
[Parallel(n_jobs=-2)]: Using backend LokyBackend with 3 concurrent workers.
[Parallel(n_jobs=-2)]: Done 250 tasks      | elapsed:    2.2s
[Parallel(n_jobs=-2)]: Done 1450 tasks    | elapsed:   12.9s
Best parameters: {'C': 0.8, 'coef0': 0.0, 'degree': 3, 'gamma': 'scale',
'kernel': 'poly'}, with F1 score of 0.80
[Parallel(n_jobs=-2)]: Done 1920 out of 1920 | elapsed:   17.5s finished
```

```
[29]: evaluate(svm_clf, test_X, test_y)
roc(svm_clf, test_X, test_y)
prc(svm_clf, test_X, test_y)
confusion(svm_clf, test_X, test_y)
plt.show()
```

```
RMSE: 0.4508
Accuracy: 0.796 ± 0.224
F1 Score: 0.78
```





## 0.7 7. Deep Neural Network

Deep neural networks (DNNs) are artificial neural (ANNs) networks with several hidden layers. Each layer is a fixed number of artificial neurons, which accept an input, process it, and send it to the next layer. The layers are organized followingly:

input layer  $\rightarrow$  hidden layers  $\rightarrow$  output layer.

Each layer has an activation function, whose choice greatly influences the overall performance. In classification tasks, the output layer yields the final class labels. We will use `Sequential model` from `Keras` as our DNN.

For the activation functions we will stick with ReLU or Rectified Linear Units. These are nearly linear functions commonly used in DNNs and provide the best results. Leaky ReLU may be used as well.

We will be choosing either Adam, Stochastic Gradient Descent or RMSProp optimizer.

Batch size will remain constant 32 and epochs 10-15, since these numbers offer the best results for the time spent learning.

We have experimented with dropout a bit, but found little to no difference when using it, so it will be kept at 0

```
[30]: import tensorflow as tf
import itertools
import itertools
```

```

import gc

import keras.backend as K
from keras.optimizers import Adam, SGD, RMSprop
from keras.models import Sequential
from keras.callbacks import EarlyStopping
from keras.layers import Dense, Dropout
from sklearn.metrics import classification_report, confusion_matrix
from keras.wrappers.scikit_learn import KerasClassifier

```

Using TensorFlow backend.

```

[31]: layer_sizes = [[[8, 16, 32, 64] for _ in range(size)] for size in range(6, 7)]
layer_combinations = list(itertools.chain.from_iterable(map(lambda sublist: list(itertools.product(*sublist)), layer_sizes)))
del layer_sizes

gc.collect()
gc.enable()

```

```

[32]: def build_net(optim, layers, lr, dropout, **kwargs):
    K.clear_session()
    model = Sequential()
    model.add(Dense(layers[0], input_shape=(train_X.shape[1],),
    ↪activation='relu'))
    model.add(Dropout(dropout))
    for layer in layers[1:]:
        model.add(Dense(layer, activation='relu'))
        model.add(Dropout(dropout))
    model.add(Dense(2, activation='softmax'))
    model.compile(loss='categorical_crossentropy',
    ↪optimizer=optim(learning_rate=lr), metrics=['accuracy'])
    return model

```

```
net_clf = KerasClassifier(build_fn=build_net, verbose=0)
```

```
layer_sizes = [[[32, 64, 128] for _ in range(size)] for size in range(3, 5)]
```

```
layer_combinations = list(itertools.chain.from_iterable(map(lambda sublist: list(itertools.pro
```

```
net_values = {"optim": [Adam, SGD, RMSprop], "epochs": [10], "batch_size": [32], "layers": lay
```

```
es = EarlyStopping(monitor='loss', min_delta=0, patience=2, verbose=0, mode='auto')
```

```
dnn_clf = get_gscv(net_clf, net_values, callbacks=[es])
```

Fitting 3 folds for each of 2880 candidates, totalling 8640 fits

1. Best parameters: {'batch\_size': 32, 'dropout': 0.1, 'epochs': 10, 'layers': (32, 64, 32, 64), 'lr': 0.004, 'optim': <class 'keras.optimizers.RMSprop'>}, with F1 score of 0.78

2. Best parameters: {'batch\_size': 32, 'dropout': 0.0, 'epochs': 15, 'layers': (8, 32, 8, 32), 'lr': 0.001, 'optim': <class 'keras.optimizers.RMSprop'>}, with F1 score of 0.77

3. Best parameters: {'batch\_size': 32, 'dropout': 0.1, 'epochs': 12, 'layers': (64, 64, 64, 32, 16), 'lr': 0.0003, 'optim': <class 'keras.optimizers.Adam'>}, with F1 score of 0.85

```
[33]: ohe = OneHotEncoder()
```

```
nn_train_y = ohe.fit_transform(y_train.to_numpy().reshape(-1, 1))
nn_test_y = ohe.transform(y_test.to_numpy().reshape(-1, 1))
```

```
[34]: # best_args = {'batch_size': 32, 'dropout': 0.1, 'epochs': 15, 'layers': (32,
→32, 8), 'lr': 0.0004, 'optim': RMSprop}
# best_args = {'batch_size': 32, 'dropout': 0.1, 'epochs': 12, 'layers': (64,
→64, 64, 64, 32, 16), 'lr': 0.0003, 'optim': Adam}
# best_args = {'batch_size': 32, 'dropout': 0.0, 'epochs': 15, 'layers': (32,
→32, 16, 8), 'lr': 0.0003, 'optim': Adam}
# best_args = {'batch_size': 32, 'dropout': 0.0, 'epochs': 12, 'layers': (64,
→64, 64, 16, 32, 64), 'lr': 0.0003, 'optim': Adam}
best_args = {'batch_size': 32, 'dropout': 0.0, 'epochs': 20, 'layers': (256,
→128, 64, 64, 32), 'lr': 0.0003, 'optim': Adam}

dnn_clf = KerasClassifier(build_fn=build_net, verbose=0, **best_args)

data = dnn_clf.fit(train_X, train_y, epochs=15, batch_size=32)
dnn_clf.model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	5376
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0

dense_4 (Dense)	(None, 64)	4160
-----		
dropout_4 (Dropout)	(None, 64)	0
-----		
dense_5 (Dense)	(None, 32)	2080
-----		
dropout_5 (Dropout)	(None, 32)	0
-----		
dense_6 (Dense)	(None, 2)	66
=====		
Total params: 52,834		
Trainable params: 52,834		
Non-trainable params: 0		
-----		

```
[35]: from sklearn.metrics import roc_curve

true_y_labels = np.argmax(nn_test_y, axis=1)

predicted_y = dnn_clf.predict(test_X)

roc_curve(true_y_labels, predicted_y)
fpr, tpr, thresholds = roc_curve(true_y_labels, predicted_y)

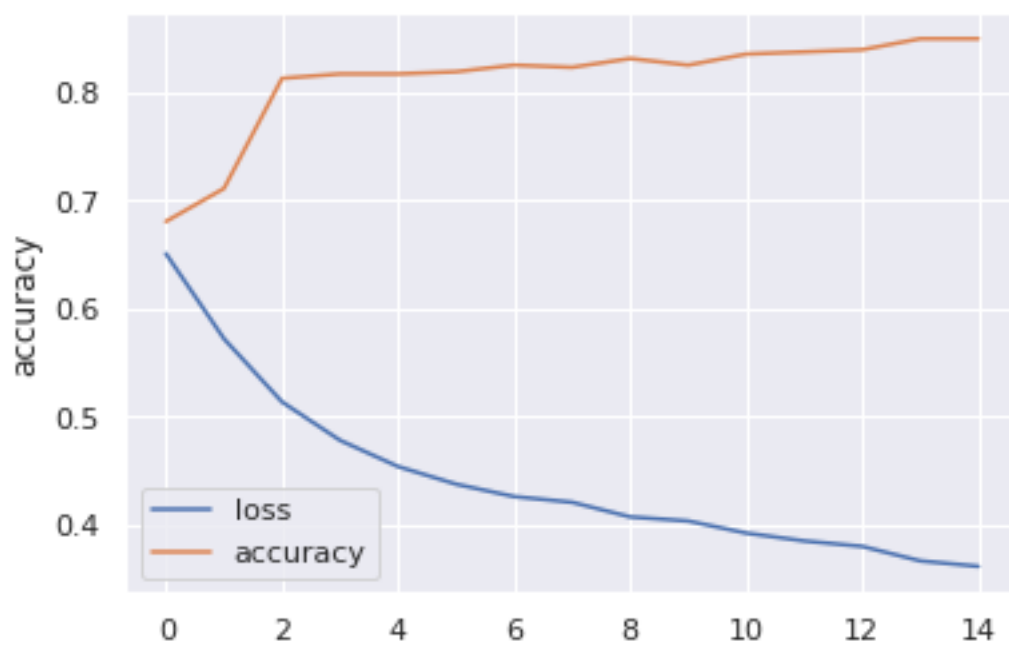
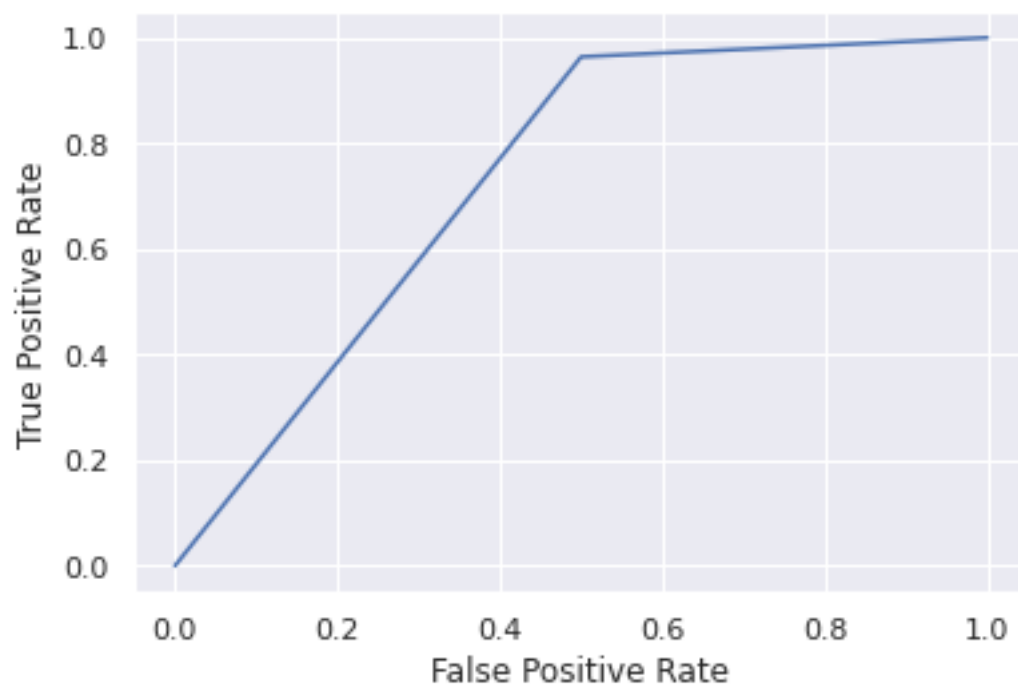
roc_data = pd.DataFrame({"False Positive Rate": fpr, "True Positive Rate": tpr})

sns.lineplot(x="False Positive Rate", y="True Positive Rate", data=roc_data)
plt.show()

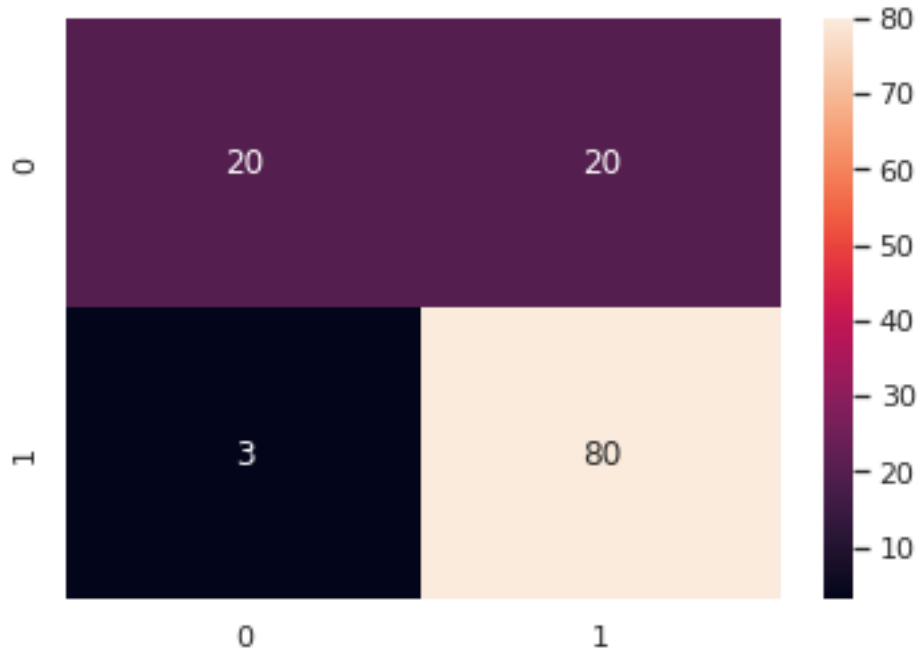
history_df = pd.DataFrame(data=data.history, columns=data.history.keys())
sns.lineplot(legend='full', y=history_df['loss'], x=range(len(data.
    ↳history['loss'])), label='loss')
sns.lineplot(legend='full', y=history_df['accuracy'], x=range(len(data.
    ↳history['accuracy'])), label='accuracy')
plt.show()

evaluate(dnn_clf, test_X, test_y)
sns.heatmap(confusion_matrix(true_y_labels, predicted_y), annot=True)
plt.show()

print('\nClassification Report')
target_names = ["Y", "N"]
print(classification_report(true_y_labels, predicted_y,
    ↳target_names=target_names))
```



RMSE: 0.4324  
Accuracy: 0.788  $\pm$  0.228  
F1 Score: 0.80



## 0.8 8. Evaluation

It is easy to see, that we have moved far beyond the performance of the baseline model. Therefore, we could assume our project reached its goal.

From the evaluation metrics it seems, that *deep neural network*, *KNN*, and *SVM* performed nearly equally. Their accuracy exceeded 80%. This is quite surprising as *KNN* can be considered as the most simple from all 4 models and yet it kept pace with them. On the other side of the spectrum is *decision tree classifier*, which had the worst evaluation metrics from all 4 models.

And finally, the winner's podium:

1. Deep Neural Network, KNN Classifier, Support Vector Machine
2. Decision Tree Classifier
3. Dummy classifier

Though keep in mind that with such small dataset the performance of all models may be influenced by the random state quite a bit.

The performance of certain models could be improved (such as the DNN classifier) by using weighted samples as the dataset is quite imbalanced.

The accuracy has quite a large differences between positive and negative samples. This is most likely caused by the fact that sampling for model training is not stratified whereas accuracy is computed using stratified cross validation.

All models seemed to have problems with recognising true negatives. In all cases the number of false negatives was greater than the number of true negatives. Though this could be due to imbalanced dataset as well.



## **0.9 9. Conclusion**

We have explored and preprocessed the dataset. From the computational side, training of the models and tuning of their hyperparameters did not take too long, in average about 35sec per model, with neural network being an exception, as it was trained with several epochs for each parameter search. Even though the dataset did not offer many records, we can conclude that the models performed overall quite well.